

```
import numpy as np
import pandas as pd
import scipy.stats
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from statsmodels.formula.api import ols
from statsmodels.stats.multicomp import pairwise_tukeyhsd
from scipy.stats import chi2_contingency
```

```
from google.colab import files
uploaded = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed

Saving Heart_disease_UCI_EDA.csv to Heart_disease_UCI_EDA (1).csv

```
df = pd.read_csv("Heart_disease_UCI_EDA.csv")
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

```
df.isna().sum()
```

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
```

```

exang      0
oldpeak    0
slope      0
ca         0
thal       0
target     0
dtype: int64

```

```
df.describe()
```

	age	sex	cp	trestbps	chol	fbs	restecg
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.528053	131.280528	245.134488	0.148515	0.514851
std	9.082101	0.466011	0.500038	16.582241	47.552910	0.356198	0.500601
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000
75%	61.000000	1.000000	1.000000	140.000000	274.500000	0.000000	1.000000
max	77.000000	1.000000	1.000000	170.000000	369.750000	1.000000	1.000000

```
df.shape
```

```
(303, 14)
```

```
df1= df[:]
```

```
df1.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null   int64
 1   sex         303 non-null   int64
 2   cp          303 non-null   int64
 3   trestbps    303 non-null   int64
 4   chol        303 non-null   float64
 5   fbs         303 non-null   int64
 6   restecg     303 non-null   int64
 7   thalach     303 non-null   float64
 8   exang       303 non-null   int64
 9   oldpeak     303 non-null   float64
10   slope       303 non-null   int64
11   ca          303 non-null   int64
12   thal        303 non-null   int64
13   target      303 non-null   int64

```

```
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

```
#Creating Dummies of Categorical Variables and dropping 1st dummy variable
```

```
Catg_vars =['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']
```

```
for i in Catg_vars:
    Catg_list = 'var'+ '_' +i
    Catg_list = pd.get_dummies(df[i], drop_first=True, prefix = i)
    df1 = df1.join(Catg_list)
df1
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 22 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
```

```

10  slope      303 non-null    int64
11  ca         303 non-null    int64
12  thal       303 non-null    int64
13  target     303 non-null    int64
14  sex_1      303 non-null    uint8
15  cp_1       303 non-null    uint8
16  fbs_1      303 non-null    uint8
17  restecg_1  303 non-null    uint8
18  exang_1    303 non-null    uint8
19  slope_2    303 non-null    uint8
20  ca_1       303 non-null    uint8
21  thal_1     303 non-null    uint8
dtypes: float64(3), int64(11), uint8(8)
memory usage: 35.6 KB

```

```

#After Creating dummies and dropping 1st dummy now drop original variable
Catg_vars = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']

```

```
df_vars = df1.columns.values.tolist()
```

```

to_keep = [i for i in df_vars if i not in Catg_vars]
# keep only those which are not in the list of data_vars

```

```
df_final = df1[to_keep]
```

```
df_final.columns.values
```

```

array(['age', 'trestbps', 'chol', 'thalach', 'oldpeak', 'target', 'sex_1',
       'cp_1', 'fbs_1', 'restecg_1', 'exang_1', 'slope_2', 'ca_1',
       'thal_1'], dtype=object)

```

```
df_final.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null    int64
 1   trestbps    303 non-null    int64
 2   chol        303 non-null    float64
 3   thalach     303 non-null    float64
 4   oldpeak     303 non-null    float64
 5   target      303 non-null    int64
 6   sex_1       303 non-null    uint8
 7   cp_1        303 non-null    uint8
 8   fbs_1       303 non-null    uint8
 9   restecg_1  303 non-null    uint8
10   exang_1     303 non-null    uint8
11   slope_2    303 non-null    uint8
12   ca_1       303 non-null    uint8
13   thal_1     303 non-null    uint8
dtypes: float64(3), int64(3), uint8(8)
memory usage: 16.7 KB

```

```
X = df_final.loc[:, df_final.columns!= 'target']
y = df_final.loc[:, df_final.columns== 'target']
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   trestbps    303 non-null   int64
2   chol       303 non-null   float64
3   thalach     303 non-null   float64
4   oldpeak     303 non-null   float64
5   sex_1       303 non-null   uint8
6   cp_1        303 non-null   uint8
7   fbs_1       303 non-null   uint8
8   restecg_1   303 non-null   uint8
9   exang_1     303 non-null   uint8
10  slope_2     303 non-null   uint8
11  ca_1        303 non-null   uint8
12  thal_1      303 non-null   uint8
dtypes: float64(3), int64(2), uint8(8)
memory usage: 14.3 KB
```

```
y
```

	target
0	1
1	1
2	1
3	1
4	1
...	...
298	0
299	0
300	0
301	0
302	0

```
303 rows × 1 columns
```

```
''' Splitting the data into Train & Test (70-30 respectively) '''
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)

train = X_train.join(y_train)

train.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 212 entries, 137 to 172
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         212 non-null    int64
 1   trestbps    212 non-null    int64
 2   chol        212 non-null    float64
 3   thalach     212 non-null    float64
 4   oldpeak     212 non-null    float64
 5   sex_1       212 non-null    uint8
 6   cp_1        212 non-null    uint8
 7   fbs_1       212 non-null    uint8
 8   restecg_1   212 non-null    uint8
 9   exang_1     212 non-null    uint8
10   slope_2     212 non-null    uint8
11   ca_1        212 non-null    uint8
12   thal_1      212 non-null    uint8
13   target      212 non-null    int64
dtypes: float64(3), int64(3), uint8(8)
memory usage: 23.2 KB

no_disease = train[train.target == 0]
len(no_disease)

94

yes_disease = train[train.target == 1]
len(yes_disease)

118

from sklearn.utils import resample

# Smote is done - over sampling
no_disease_os = resample(no_disease,
                        replace = True,
                        n_samples = len(yes_disease),
                        random_state = 14)

train_os = pd.concat([yes_disease, no_disease_os])

train_os.target.value_counts()

```

```
1    118
0    118
Name: target, dtype: int64
```

```
X_train_os = train_os.loc[:, train_os.columns != 'target']
y_train_os = train_os.loc[:, train_os.columns == 'target']
```

#Recursive Feature Elimination

```
from sklearn import datasets
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
```

```
logreg = LogisticRegression(max_iter=10000000)
```

```
rfe = RFE(logreg, n_features_to_select=11)
```

```
rfe = rfe.fit(X_train_os, y_train_os.values.ravel())
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:354: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:354: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:354: ConvergenceWarning:
ConvergenceWarning,
```



```
rfe.n_features_to_select
```

```
11
```

```
X_train_os.columns[rfe.get_support()]
```

```
Index(['trestbps', 'thalach', 'oldpeak', 'sex_1', 'cp_1', 'fbs_1', 'restecg_1',
      'exang_1', 'slope_2', 'ca_1', 'thal_1'],
      dtype='object')
```

```
cols = X_train_os.columns[rfe.get_support()]
```

```
cols.to_list()
```

```
['trestbps',
 'thalach',
 'oldpeak',
 'sex_1',
 'cp_1',
 'fbs_1',
 'restecg_1',
```

```
'exang_1',
'slope_2',
'ca_1',
'thal_1']
```

#Logistic Model by statistic apporach

```
#_____sm model to see p_values
```

```
x1 = X_train_os[cols]
```

```
x1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 236 entries, 137 to 186
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   trestbps    236 non-null    int64
1   thalach     236 non-null    float64
2   oldpeak     236 non-null    float64
3   sex_1       236 non-null    uint8
4   cp_1        236 non-null    uint8
5   fbs_1       236 non-null    uint8
6   restecg_1   236 non-null    uint8
7   exang_1     236 non-null    uint8
8   slope_2     236 non-null    uint8
9   ca_1        236 non-null    uint8
10  thal_1      236 non-null    uint8
dtypes: float64(2), int64(1), uint8(8)
memory usage: 9.2 KB
```

```
y_train_os.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 236 entries, 137 to 186
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   target  236 non-null    int64
dtypes: int64(1)
memory usage: 3.7 KB
```

```
y_train_os.value_counts()
```

```
target
1      118
0      118
dtype: int64
```

```
y1 = y_train_os
```

```
y1
```


	target
137	1
106	1
44	1
139	1
156	1
...	...
263	0
257	0
221	0
221	0
186	0

236 rows × 1 columns

```
#_____Stats model
import statsmodels.api as sm
```

```
x1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 236 entries, 137 to 186
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   trestbps    236 non-null    int64
1   thalach     236 non-null    float64
2   oldpeak     236 non-null    float64
3   sex_1       236 non-null    uint8
4   cp_1        236 non-null    uint8
5   fbs_1       236 non-null    uint8
6   restecg_1   236 non-null    uint8
7   exang_1     236 non-null    uint8
8   slope_2     236 non-null    uint8
9   ca_1        236 non-null    uint8
10  thal_1      236 non-null    uint8
dtypes: float64(2), int64(1), uint8(8)
memory usage: 9.2 KB
```

```
y1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 236 entries, 137 to 186
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   target  236 non-null    int64
```

```
dtypes: int64(1)
memory usage: 3.7 KB
```

```
logit_model = sm.Logit(y1,x1)
```

```
result = logit_model.fit(method='bfgs')
```

Warning: Maximum number of iterations has been exceeded.

Current function value: 0.290243

Iterations: 35

Function evaluations: 39

Gradient evaluations: 39

/usr/local/lib/python3.7/dist-packages/statsmodels/base/model.py:512: ConvergenceWarning
"Check mle_retvals", ConvergenceWarning)

```
print(result.summary2())
```

Results: Logit

```
=====
Model:                Logit                Pseudo R-squared: 0.581
Dependent Variable: target                AIC:                158.9947
Date:                2021-11-27 07:59 BIC:                197.0969
No. Observations:    236                Log-Likelihood:    -68.497
Df Model:            10                LL-Null:          -163.58
Df Residuals:        225                LLR p-value:      1.8016e-35
Converged:            0.0000                Scale:          1.0000
=====
```

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
trestbps	-0.0040	0.0099	-0.4054	0.6852	-0.0233	0.0153
thalach	0.0197	0.0089	2.2180	0.0266	0.0023	0.0371
oldpeak	-0.5144	0.2615	-1.9674	0.0491	-1.0269	-0.0020
sex_1	-1.2426	0.5543	-2.2419	0.0250	-2.3290	-0.1563
cp_1	1.7922	0.5053	3.5467	0.0004	0.8018	2.7825
fbs_1	-0.6882	0.6324	-1.0883	0.2765	-1.9277	0.5513
restecg_1	-0.5831	0.4842	-1.2043	0.2285	-1.5321	0.3659
exang_1	-0.7763	0.4994	-1.5545	0.1201	-1.7551	0.2025
slope_2	0.4471	0.5568	0.8030	0.4220	-0.6442	1.5383
ca_1	-1.8748	0.4714	-3.9773	0.0001	-2.7987	-0.9509
thal_1	-1.1527	0.5104	-2.2587	0.0239	-2.1530	-0.1524

```
=====
```

#Logistic model by SK learn method

```
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
```

```
logreg= LogisticRegression(solver= 'sag')
logreg.fit(x1, y1)
```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:985: DataConversionWarning
y = column_or_1d(y, warn=True)

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:354: ConvergenceWarning,
LogisticRegression(solver='sag')
```

```
## X_test should also have only 2 columns
X_test2= X_test[cols]
```

```
X_test2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 91 entries, 225 to 238
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   trestbps    91 non-null    int64
1   thalach     91 non-null    float64
2   oldpeak     91 non-null    float64
3   sex_1       91 non-null    uint8
4   cp_1        91 non-null    uint8
5   fbs_1       91 non-null    uint8
6   restecg_1   91 non-null    uint8
7   exang_1     91 non-null    uint8
8   slope_2     91 non-null    uint8
9   ca_1        91 non-null    uint8
10  thal_1      91 non-null    uint8
dtypes: float64(2), int64(1), uint8(8)
memory usage: 3.6 KB
```

```
y_pred= logreg.predict(X_test2)
```

```
log_score= logreg.score(X_test2, y_test)
```

```
print("Accuracy of logistic regression classifier on test data:{}".format(log_score))
```

```
Accuracy of logistic regression classifier on test data:0.7032967032967034
```

```
from sklearn.metrics import confusion_matrix
```

```
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)
```

```
[[28 16]
 [11 36]]
```

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support

0           0.72     0.64     0.67         44
```

	1	0.69	0.77	0.73	47
accuracy				0.70	91
macro avg		0.71	0.70	0.70	91
weighted avg		0.70	0.70	0.70	91


```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve

logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test2))
logit_roc_auc
```

```
0.7011605415860734
```

```
""" Area under curve is 0.701 """
```

```
fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test2)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([-0.1, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate (1-specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.title('ROC - AUC for Logistic Model')
plt.legend(loc="lower right")
plt.show()
```



Decision Tree Model

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
```

```

#    Column    Non-Null Count  Dtype
---  -
0    age       303 non-null    int64
1    sex       303 non-null    int64
2    cp        303 non-null    int64
3    trestbps  303 non-null    int64
4    chol      303 non-null    float64
5    fbs       303 non-null    int64
6    restecg   303 non-null    int64
7    thalach   303 non-null    float64
8    exang     303 non-null    int64
9    oldpeak   303 non-null    float64
10   slope     303 non-null    int64
11   ca        303 non-null    int64
12   thal      303 non-null    int64
13   target    303 non-null    int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```
df2=df[:]
```

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#    Column    Non-Null Count  Dtype
---  -
0    age       303 non-null    int64
1    sex       303 non-null    int64
2    cp        303 non-null    int64
3    trestbps  303 non-null    int64
4    chol      303 non-null    float64
5    fbs       303 non-null    int64
6    restecg   303 non-null    int64
7    thalach   303 non-null    float64
8    exang     303 non-null    int64
9    oldpeak   303 non-null    float64
10   slope     303 non-null    int64
11   ca        303 non-null    int64
12   thal      303 non-null    int64
13   target    303 non-null    int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

from sklearn.preprocessing import LabelEncoder
LE= LabelEncoder()

```

```

df2['sex']= LE.fit_transform(df2['sex'])
df2['cp']= LE.fit_transform(df2['cp'])
df2['fbs']= LE.fit_transform(df2['fbs'])
df2['restecg']= LE.fit_transform(df2['restecg'])
df2['exang']= LE.fit_transform(df2['exang'])
df2['slope']= LE.fit_transform(df2['slope'])
df2['ca']= LE.fit_transform(df2['ca'])

```

```
df2['thal'] = LE.fit_transform(df2['thal'])
df2
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
0	63	1	1	145	233.0	1	0	150.0	0	2.3	0	0
1	37	1	1	130	250.0	0	1	187.0	0	3.5	0	0
2	41	0	1	130	204.0	0	0	172.0	0	1.4	1	0
3	56	1	1	120	236.0	0	1	178.0	0	0.8	1	0
4	57	0	0	120	354.0	0	1	163.0	1	0.6	1	0
...
298	57	0	0	140	241.0	0	1	123.0	1	0.2	0	0
299	45	1	1	110	264.0	0	1	132.0	0	1.2	0	0
300	68	1	0	144	193.0	1	1	141.0	0	3.4	0	1
301	57	1	0	130	131.0	0	1	115.0	1	1.2	0	1
302	57	0	1	130	236.0	0	0	174.0	0	0.0	0	1

303 rows × 14 columns

```
X = df2.loc[:, df2.columns != 'target']
y = df2.loc[:, df2.columns == 'target']
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    float64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    float64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.9 KB
```

```
y
```

	target
0	1
1	1
2	1
3	1
4	1
...	...
298	0
299	0
300	0
301	0
302	0

```
'''Fit Tree'''
```

```
#train test - split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
#fit tree on train data
```

```
#model
```

```
clf = DecisionTreeClassifier()
```

```
#Fit Classifier model on train set
```

```
clf.fit(X_train, y_train)
```

```
DecisionTreeClassifier()
```

```
#Predict/estimate_train X_train
```

```
y_pred_train = clf.predict(X_train)
```

```
y_pred_train
```

```
array([1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0,
       0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1,
       0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0,
       0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0,
       0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1,
       0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
       0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0,
       0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0,
```

```
1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,  
1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0])
```

```
#Predict/estimate_test X_test  
y_pred_test = clf.predict(X_test)
```

```
y_pred_test
```

```
array([0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0,  
       0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0,  
       0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
       1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1,  
       0, 1, 0])
```

```
#See the train  
from sklearn import tree
```

```
tree.plot_tree(clf.fit(X_train, y_train))
```



```
[Text(180.20117647058825, 206.56799999999998, 'X[12] <= 0.5\ngini = 0.494\nsamples =
Text(110.28705882352942, 184.824, 'X[9] <= 1.7\ngini = 0.303\nsamples = 118\nvalue =
Text(78.7764705882353, 163.07999999999998, 'X[11] <= 0.5\ngini = 0.222\nsamples = 16
Text(47.265882352941176, 141.336, 'X[0] <= 58.5\ngini = 0.139\nsamples = 80\nvalue =
Text(31.51058823529412, 119.592, 'X[7] <= 156.5\ngini = 0.064\nsamples = 60\nvalue =
Text(23.632941176470588, 97.848, 'X[3] <= 129.0\ngini = 0.219\nsamples = 16\nvalue =
Text(15.75529411764706, 76.10399999999998, 'X[1] <= 0.5\ngini = 0.48\nsamples = 5\n
Text(7.87764705882353, 54.360000000000014, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]
Text(23.632941176470588, 54.360000000000014, 'gini = 0.0\nsamples = 2\nvalue = [2, 6
Text(31.51058823529412, 76.10399999999998, 'gini = 0.0\nsamples = 11\nvalue = [0, 11
Text(39.38823529411765, 97.848, 'gini = 0.0\nsamples = 44\nvalue = [0, 44]'),
Text(63.02117647058824, 119.592, 'X[4] <= 243.5\ngini = 0.32\nsamples = 20\nvalue =
Text(55.14352941176471, 97.848, 'gini = 0.0\nsamples = 9\nvalue = [0, 9]'),
Text(70.89882352941177, 97.848, 'X[0] <= 64.5\ngini = 0.463\nsamples = 11\nvalue = [
Text(63.02117647058824, 76.10399999999998, 'X[7] <= 170.5\ngini = 0.5\nsamples = 8\nr
Text(55.14352941176471, 54.360000000000014, 'X[3] <= 165.0\ngini = 0.444\nsamples =
Text(47.265882352941176, 32.615999999999985, 'X[4] <= 352.375\ngini = 0.32\nsamples
Text(39.38823529411765, 10.872000000000014, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]
Text(55.14352941176471, 10.872000000000014, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]
Text(63.02117647058824, 32.615999999999985, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]
Text(70.89882352941177, 54.360000000000014, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]
Text(78.7764705882353, 76.10399999999998, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]')
Text(110.28705882352942, 141.336, 'X[2] <= 0.5\ngini = 0.434\nsamples = 22\nvalue =
Text(94.53176470588235, 119.592, 'X[1] <= 0.5\ngini = 0.32\nsamples = 5\nvalue = [4,
Text(86.65411764705883, 97.848, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(102.40941176470588, 97.848, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(126.04235294117647, 119.592, 'X[4] <= 172.0\ngini = 0.291\nsamples = 17\nvalue
Text(118.16470588235295, 97.848, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(133.92000000000002, 97.848, 'X[4] <= 279.5\ngini = 0.219\nsamples = 16\nvalue =
Text(126.04235294117647, 76.10399999999998, 'gini = 0.0\nsamples = 9\nvalue = [0, 9]
Text(141.79764705882354, 76.10399999999998, 'X[4] <= 285.0\ngini = 0.408\nsamples =
Text(133.92000000000002, 54.360000000000014, 'gini = 0.0\nsamples = 1\nvalue = [1, 6
Text(149.67529411764707, 54.360000000000014, 'X[7] <= 154.5\ngini = 0.278\nsamples =
Text(141.79764705882354, 32.615999999999985, 'X[0] <= 54.5\ngini = 0.5\nsamples = 2\
Text(133.92000000000002, 10.872000000000014, 'gini = 0.0\nsamples = 1\nvalue = [0, 1
Text(149.67529411764707, 10.872000000000014, 'gini = 0.0\nsamples = 1\nvalue = [1, 6
Text(157.5529411764706, 32.615999999999985, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]
Text(141.79764705882354, 163.07999999999998, 'X[0] <= 45.5\ngini = 0.492\nsamples =
Text(133.92000000000002, 141.336, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(149.67529411764707, 141.336, 'X[0] <= 63.5\ngini = 0.426\nsamples = 13\nvalue =
Text(141.79764705882354, 119.592, 'gini = 0.0\nsamples = 8\nvalue = [8, 0]'),
Text(157.5529411764706, 119.592, 'X[7] <= 111.5\ngini = 0.32\nsamples = 5\nvalue = [
Text(149.67529411764707, 97.848, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(165.43058823529412, 97.848, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),
Text(250.11529411764707, 184.824, 'X[2] <= 0.5\ngini = 0.359\nsamples = 94\nvalue =
Text(208.75764705882352, 163.07999999999998, 'X[9] <= 0.5\ngini = 0.19\nsamples = 66
Text(181.18588235294118, 141.336, 'X[7] <= 117.0\ngini = 0.43\nsamples = 16\nvalue =
Text(173.30823529411765, 119.592, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(189.0635294117647, 119.592, 'X[4] <= 217.0\ngini = 0.337\nsamples = 14\nvalue =
Text(181.18588235294118, 97.848, 'X[11] <= 0.5\ngini = 0.5\nsamples = 6\nvalue = [3,
Text(173.30823529411765, 76.10399999999998, 'X[6] <= 0.5\ngini = 0.375\nsamples = 4\
Text(165.43058823529412, 54.360000000000014, 'gini = 0.0\nsamples = 1\nvalue = [1, 6
Text(181.18588235294118, 54.360000000000014, 'gini = 0.0\nsamples = 3\nvalue = [0, 3
Text(189.0635294117647, 76.10399999999998, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]
Text(196.94117647058823, 97.848, 'gini = 0.0\nsamples = 8\nvalue = [8, 0]'),
Text(236.3294117647059, 141.336, 'X[0] <= 65.5\ngini = 0.077\nsamples = 50\nvalue =
Text(220.57411764705884, 119.592, 'X[3] <= 112.0\ngini = 0.042\nsamples = 47\nvalue
```

```
#Model has learnt unnecessary things
#Need to optimize
```

```
from sklearn.metrics import accuracy_score
```

```
print(round(accuracy_score(y_train,y_pred_train), 2))
```

```
1.0
```

```
print(round(accuracy_score(y_test,y_pred_test), 2))
```

```
0.66
```

```
# Accuracy of train data is 1
# Accuracy of test data is 0.66
```

```
from sklearn import tree
```

```
path = clf.cost_complexity_pruning_path(X_train, y_train)
```

```
path
```

```
{'ccp_alphas': array([0.          , 0.00393082, 0.00408805, 0.0043239 , 0.0045283 ,
 0.00581761, 0.00628931, 0.0067086 , 0.006798 , 0.00707547,
 0.00754717, 0.00754717, 0.00808625, 0.00825472, 0.00934166,
 0.01027254, 0.01189882, 0.01479953, 0.01782345, 0.02469498,
 0.03423776, 0.16576779]),
'impurities': array([0.          , 0.00786164, 0.02421384, 0.03286164, 0.05097484,
 0.10333333, 0.10962264, 0.11633124, 0.12312924, 0.13728018,
 0.14482735, 0.15237452, 0.16046078, 0.16871549, 0.18739882,
 0.19767136, 0.22146899, 0.25106805, 0.2688915 , 0.29358648,
 0.32782424, 0.49359203])}
```

```
alphas = path['ccp_alphas']
```

```
alphas
```

```
array([0.          , 0.00393082, 0.00408805, 0.0043239 , 0.0045283 ,
 0.00581761, 0.00628931, 0.0067086 , 0.006798 , 0.00707547,
 0.00754717, 0.00754717, 0.00808625, 0.00825472, 0.00934166,
 0.01027254, 0.01189882, 0.01479953, 0.01782345, 0.02469498,
 0.03423776, 0.16576779])
```

```
acrcy_train, acrcy_test = [],[]
```

```
for i in alphas:
    clf = DecisionTreeClassifier(ccp_alpha=i)
```

```
    clf.fit(X_train, y_train)
```

```
y_pred_train = clf.predict(X_train)
y_pred_test = clf.predict(X_test)

acrcy_train.append(accuracy_score(y_train, y_pred_train))
acrcy_test.append(accuracy_score(y_test,y_pred_test))
```

acrcy_train


```
[1.0,
 0.9952830188679245,
 0.9858490566037735,
 0.9811320754716981,
 0.9716981132075472,
 0.9433962264150944,
 0.9386792452830188,
 0.9339622641509434,
 0.9292452830188679,
 0.9150943396226415,
 0.9056603773584906,
 0.9056603773584906,
 0.9056603773584906,
 0.9009433962264151,
 0.8915094339622641,
 0.8820754716981132,
 0.8679245283018868,
 0.839622641509434,
 0.8113207547169812,
 0.8018867924528302,
 0.7924528301886793,
 0.5566037735849056]
```

acrcy_test

```
[0.7252747252747253,
 0.6813186813186813,
 0.7802197802197802,
 0.7582417582417582,
 0.8131868131868132,
 0.7912087912087912,
 0.8021978021978022,
 0.8131868131868132,
 0.7802197802197802,
 0.7802197802197802,
 0.7912087912087912,
 0.8021978021978022,
 0.8021978021978022,
 0.8021978021978022,
 0.8021978021978022,
 0.8021978021978022,
 0.7472527472527473,
 0.7472527472527473,
 0.7252747252747253,
 0.7142857142857143,
 0.7032967032967034,
 0.5164835164835165]
```

```
# now we have scores
# lets, plot
```

```
sns.set()
plt.figure(figsize = (14,7))
sns.lineplot(y =acrcy_train, x = alphas, label = 'Train_Accuracy')
sns.lineplot(y =acrcy_test, x = alphas, label = 'Test_Accuracy')
plt.xticks(ticks=np.arange(0.00,0.2,0.01))
plt.xlabel('alphas')
plt.ylabel('accuracy')
plt.show()
```



```
# _____ with ccp = 0.015
clf = DecisionTreeClassifier(ccp_alpha=0.015, random_state = 14)
```

```
clf.fit(X_train,y_train)
```

```
DecisionTreeClassifier(ccp_alpha=0.015, random_state=14)
```

```
y_pred_train = clf.predict(X_train)
```

```
y_pred_test = clf.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
print(round(accuracy_score(y_train,y_pred_train), 2))
```

```
0.84
```

```
print(round(accuracy_score(y_test,y_pred_test), 2))
```

```
0.75
```

```
### Confusion Matrix
```

```
from sklearn.metrics import confusion_matrix
```

```
from sklearn.metrics import classification_report
```

```
confusion_matrix = confusion_matrix(y_test, y_pred_test)
```

```
print(confusion_matrix)
```

```
[[27 17]
 [ 6 41]]
```

```
### Classification Report
```

```
from sklearn.metrics import classification_report
```

```
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.82	0.61	0.70	44
1	0.71	0.87	0.78	47
accuracy			0.75	91
macro avg	0.76	0.74	0.74	91
weighted avg	0.76	0.75	0.74	91

```
##### ROC AUC Curve
```

```
from sklearn.metrics import roc_auc_score
```

```
from sklearn.metrics import roc_curve
```

```
from sklearn.metrics import roc_curve, auc, roc_auc_score
```

```
predictedProbability = clf.predict_proba(X_test)[: , 1]
```

```
fpr, tpr, thresholds = metrics.roc_curve(y_test, predictedProbability)
```

```
fpr
```

```
array([0.          , 0.27272727, 0.38636364, 0.45454545, 0.52272727,
       1.          ])
```

```
tpr
```

```
array([0.          , 0.70212766, 0.87234043, 0.91489362, 0.93617021,
       1.          ])
```


```
thresholds
```

```
array([1.87254902, 0.87254902, 0.7       , 0.4375     , 0.125       ,
       0.10606061])
```

```
dff = pd.DataFrame(dict(fpr = fpr,tpr = tpr))
auc = auc(fpr,tpr)
auc
```

```
0.7712765957446808
```

```
plt.figure()
lw = 2
plt.plot(fpr, tpr, color = 'darkorange',
         lw =lw, label = 'ROC Curve (area = %0.2f)' %auc)
plt.plot([0,1],[0,1], color='navy', lw = lw, linestyle = '--')
plt.xlim([-0.1, 1.0])
plt.ylim([0.0, 1.1])
plt.xlabel('1-Specificity(False Positive Rate)')
plt.ylabel('Sensitivity(True Positive Rate)')
plt.title("ROC - AUC for Decision Tree Model")
plt.legend(loc = "lower right")
plt.show()
```



Random Forest

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
```

```

7   thalach   303 non-null   float64
8   exang     303 non-null   int64
9   oldpeak   303 non-null   float64
10  slope     303 non-null   int64
11  ca        303 non-null   int64
12  thal      303 non-null   int64
13  target    303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```
df2=df[:]
```

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

from sklearn.preprocessing import LabelEncoder
LE= LabelEncoder()

```

```

df2['sex']= LE.fit_transform(df2['sex'])
df2['cp']= LE.fit_transform(df2['cp'])
df2['fbs']= LE.fit_transform(df2['fbs'])
df2['restecg']= LE.fit_transform(df2['restecg'])
df2['exang']= LE.fit_transform(df2['exang'])
df2['slope']= LE.fit_transform(df2['slope'])
df2['ca']= LE.fit_transform(df2['ca'])
df2['thal']= LE.fit_transform(df2['thal'])
df2

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
0	63	1	1	145	233.0	1	0	150.0	0	2.3	0	0
1	37	1	1	130	250.0	0	1	187.0	0	3.5	0	0
2	41	0	1	130	204.0	0	0	172.0	0	1.4	1	0
3	56	1	1	120	236.0	0	1	178.0	0	0.8	1	0
4	57	0	0	120	354.0	0	1	163.0	1	0.6	1	0
...
298	57	0	0	140	241.0	0	1	123.0	1	0.2	0	0
299	45	1	1	110	264.0	0	1	132.0	0	1.2	0	0
300	68	1	0	144	193.0	1	1	141.0	0	3.4	0	1
301	57	0	1	130	250.0	0	0	174.0	0	0.0	0	1

```
X = df2.loc[:, df2.columns != 'target']
```

```
y = df2.loc[:, df2.columns == 'target']
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    float64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    float64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.9 KB
```

```
y
```


	target
0	1
1	1
2	1
3	1
4	1
...	...

```
'''Fit Tree'''
```

```
#train test - split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```

```
#import the classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
#Create Classifier object
```


```
#in our previous experiment, we found ccp_alphas = 0.015 has the best accuracy
```

```
clf_rf = RandomForestClassifier(n_estimators =100, ccp_alpha= 0.015, random_state = 14)
```

```
#fit the classifier with x and y data = train
```

```
mod_rf = clf_rf.fit(X_train, y_train)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DataConversionWarning
```



```
#Prediction
```

```
y_train_pred = mod_rf.predict(X_train)
```

```
y_train_pred
```

```
array([1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0,
       0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1,
       0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0,
       0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0,
       1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1,
       0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
       0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0,
       0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0,
       1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,
       1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1])
```

```
#Prediction
```

```
y_test_pred = mod_rf.predict(X_test)
```

```
y_test_pred
```

```
array([0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0,
       0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0,
       1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1,
       1, 1, 0])
```

```
from sklearn.metrics import accuracy_score
```

```
print(round(accuracy_score(y_train,y_train_pred), 2))
```

```
0.9
```

```
print(round(accuracy_score(y_test,y_test_pred), 2))
```

```
0.81
```

```
#_____ Extract Feature Importance
fi = pd.DataFrame({'feature': list(X_train.columns),
                   'importance': mod_rf.feature_importances_}).\
    sort_values('importance', ascending = False)
```

```
fi.head()
```

	feature	importance
12	thal	0.203829
11	ca	0.156076
9	oldpeak	0.115380
2	cp	0.113929
8	exang	0.100628

```
# Accurarcy 2 cells above is 0.9 & 0.81 for Train & test (respectively)
```

```
# This accuracy is for having all columns as features in our model
```

```
# Lets build a model keeping 5 best features
```

```
# that is keeping 'cp', 'exang', 'oldpeak', 'ca', 'thal' only
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
#Create Classifier object
```

```
#in our previous experiment Decision Tree model,
```

```
#we found ccp_alphas = 0.035 has the best accuracy
```

```
clf_rf1 = RandomForestClassifier(n_estimators =100, ccp_alpha= 0.015, random_state = 14)
```

```
# fit the classifier with x and y data=TRAIN,
```

```
#this time with Failure_Type only
```

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```

Int64Index: 212 entries, 137 to 172
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         212 non-null   int64
1   sex         212 non-null   int64
2   cp          212 non-null   int64
3   trestbps    212 non-null   int64
4   chol        212 non-null   float64
5   fbs         212 non-null   int64
6   restecg     212 non-null   int64
7   thalach     212 non-null   float64
8   exang       212 non-null   int64
9   oldpeak     212 non-null   float64
10  slope       212 non-null   int64
11  ca          212 non-null   int64
12  thal        212 non-null   int64
dtypes: float64(3), int64(10)
memory usage: 23.2 KB

```

```

X_train1 = X_train.iloc[ : ,[2,8,9,11,12]]
X_train1.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 212 entries, 137 to 172
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   cp          212 non-null   int64
1   exang       212 non-null   int64
2   oldpeak     212 non-null   float64
3   ca          212 non-null   int64
4   thal        212 non-null   int64
dtypes: float64(1), int64(4)
memory usage: 9.9 KB


```

```
mod_rf1 = clf_rf1.fit(X_train1, y_train)
```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: DataConversionWarning:
  """Entry point for launching an IPython kernel.

```



```

#Prediction
y_train_pred1 = mod_rf1.predict(X_train1)
y_train_pred1

```

```

array([1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0,
       0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1,
       0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0,
       0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0,
       1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
       0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0,
       0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0,
       1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0,
       1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1])

```

```
X_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 91 entries, 225 to 238
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         91 non-null    int64
1   sex         91 non-null    int64
2   cp          91 non-null    int64
3   trestbps    91 non-null    int64
4   chol        91 non-null    float64
5   fbs         91 non-null    int64
6   restecg     91 non-null    int64
7   thalach     91 non-null    float64
8   exang       91 non-null    int64
9   oldpeak     91 non-null    float64
10  slope       91 non-null    int64
11  ca          91 non-null    int64
12  thal        91 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 10.0 KB
```

```
X_test1 = X_test.iloc[ : ,[2,8,9,11,12]]
X_test1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 91 entries, 225 to 238
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   cp          91 non-null    int64
1   exang       91 non-null    int64
2   oldpeak     91 non-null    float64
3   ca          91 non-null    int64
4   thal        91 non-null    int64
dtypes: float64(1), int64(4)
memory usage: 4.3 KB
```

```
#Prediction
```

```
y_test_pred1 = mod_rf1.predict(X_test1)
y_test_pred1
```

```
array([0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0,
       0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0,
       1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1,
       1, 1, 1])
```

```
from sklearn.metrics import accuracy_score
```

```
print(round(accuracy_score(y_train,y_train_pred1), 2))
```

0.86

```
print(round(accuracy_score(y_test,y_test_pred1), 2))
```

0.81

```
### There is no much difference in accuarcy
#Earlier train accuracy = 0.9 now with 4 features its 0.86
#Earlier test accuracy = 0.81 now with 4 features its 0.81
```

```
### Confusion Matrix
from sklearn.metrics import confusion_matrix
```

```
confusion_matrix = confusion_matrix(y_test, y_test_pred1)
print(confusion_matrix)
```

```
[[31 13]
 [ 4 43]]
```

```
### Classification Report
print(classification_report(y_test, y_test_pred1))
```

	precision	recall	f1-score	support
0	0.89	0.70	0.78	44
1	0.77	0.91	0.83	47
accuracy			0.81	91
macro avg	0.83	0.81	0.81	91
weighted avg	0.82	0.81	0.81	91

```
##### ROC AUC Curve
```

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_curve, auc, roc_auc_score
```

```
predictedProbability1 = mod_rf1.predict_proba(X_test1)[: , 1]
fpr, tpr, thresholds = metrics.roc_curve(y_test, predictedProbability1)
```

```
fpr
```

```
array([0.          , 0.          , 0.          , 0.          , 0.          ,
        0.          , 0.02272727, 0.06818182, 0.06818182, 0.09090909,
        0.09090909, 0.13636364, 0.13636364, 0.13636364, 0.15909091,
        0.15909091, 0.18181818, 0.18181818, 0.20454545, 0.20454545,
        0.22727273, 0.25          , 0.29545455, 0.29545455, 0.34090909,
        0.36363636, 0.56818182, 0.56818182, 0.63636364, 0.65909091,
        0.75          , 0.79545455, 0.84090909, 0.88636364, 0.93181818,
        0.97727273, 1.          ])
```

tpr

```
array([0.          , 0.34042553, 0.36170213, 0.40425532, 0.46808511,
       0.4893617  , 0.5106383  , 0.55319149, 0.57446809, 0.59574468,
       0.61702128, 0.63829787, 0.68085106, 0.72340426, 0.72340426,
       0.76595745, 0.76595745, 0.78723404, 0.80851064, 0.85106383,
       0.87234043, 0.91489362, 0.91489362, 0.93617021, 0.93617021,
       0.95744681, 0.95744681, 0.9787234  , 0.9787234  , 1.          ,
       1.          , 1.          , 1.          , 1.          , 1.          ,
       1.          , 1.          ])
```

thresholds

```
array([1.89541299, 0.89541299, 0.87927376, 0.87323944, 0.86817101,
       0.85796746, 0.84706323, 0.83106041, 0.82920144, 0.78756654,
       0.77792825, 0.76811059, 0.72797353, 0.71908852, 0.71394146,
       0.69817409, 0.69024108, 0.67850108, 0.66600645, 0.61190299,
       0.60839611, 0.51014379, 0.5002123  , 0.4982503  , 0.45813156,
       0.42562654, 0.28671954, 0.27201602, 0.20334771, 0.17894051,
       0.11768904, 0.11394835, 0.10070404, 0.09012524, 0.08501373,
       0.082074  , 0.081574  ])
```

```
dff1 = pd.DataFrame(dict(fpr = fpr,tpr = tpr))
```


```
auc1 = auc(fpr,tpr)
```

```
auc1
```

```
0.9008704061895552
```

```
''' Area Under Curve is 0.9 '''
```

```
plt.figure()
lw = 2
plt.plot(fpr, tpr, color = 'green',
         lw =lw, label = 'ROC Curve (area = %0.2f)' %auc1)
plt.plot([0,1],[0,1], color='navy', lw = lw, linestyle = '--')
plt.xlim([-0.1, 1.0])
plt.ylim([0.0, 1.1])
plt.xlabel('1-Specificity(False Positive Rate)')
plt.ylabel('Sensitivity(True Positive Rate)')
plt.title("ROC - AUC for Random forest")
plt.legend(loc = "lower right")
plt.show()
```



Gradient Boosting Model

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
from sklearn.model_selection import GridSearchCV
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

```
df2=df[:]
```

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
```

```

9    oldpeak    303 non-null    float64
10   slope      303 non-null    int64
11   ca         303 non-null    int64
12   thal       303 non-null    int64
13   target     303 non-null    int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

from sklearn.preprocessing import LabelEncoder
LE= LabelEncoder()

```

```

df2['sex']= LE.fit_transform(df2['sex'])
df2['cp']= LE.fit_transform(df2['cp'])
df2['fbs']= LE.fit_transform(df2['fbs'])
df2['restecg']= LE.fit_transform(df2['restecg'])
df2['exang']= LE.fit_transform(df2['exang'])
df2['slope']= LE.fit_transform(df2['slope'])
df2['ca']= LE.fit_transform(df2['ca'])
df2['thal']= LE.fit_transform(df2['thal'])
df2

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
0	63	1	1	145	233.0	1	0	150.0	0	2.3	0	0
1	37	1	1	130	250.0	0	1	187.0	0	3.5	0	0
2	41	0	1	130	204.0	0	0	172.0	0	1.4	1	0
3	56	1	1	120	236.0	0	1	178.0	0	0.8	1	0
4	57	0	0	120	354.0	0	1	163.0	1	0.6	1	0
...
298	57	0	0	140	241.0	0	1	123.0	1	0.2	0	0
299	45	1	1	110	264.0	0	1	132.0	0	1.2	0	0
300	68	1	0	144	193.0	1	1	141.0	0	3.4	0	1
301	57	1	0	130	131.0	0	1	115.0	1	1.2	0	1
302	57	0	1	130	236.0	0	0	174.0	0	0.0	0	1

303 rows × 14 columns

```

X = df2.loc[:, df2.columns != 'target']
y = df2.loc[:, df2.columns == 'target']

```

```
X.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -

```



```

0   age      303 non-null   int64
1   sex      303 non-null   int64
2   cp       303 non-null   int64
3   trestbps 303 non-null   int64
4   chol     303 non-null   float64
5   fbs      303 non-null   int64
6   restecg  303 non-null   int64
7   thalach  303 non-null   float64
8   exang    303 non-null   int64
9   oldpeak  303 non-null   float64
10  slope    303 non-null   int64
11  ca       303 non-null   int64
12  thal     303 non-null   int64

```

```
dtypes: float64(3), int64(10)
```

```
memory usage: 30.9 KB
```

y

	target
0	1
1	1
2	1
3	1
4	1
...	...
298	0
299	0
300	0
301	0
302	0

303 rows × 1 columns

```
'''Fit Tree'''
```

```
#train test - split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```

```
GB = GradientBoostingClassifier()
```

```
GB_mod = GB.fit(X_train, y_train)
```

```

/usr/local/lib/python3.7/dist-packages/sklearn/ensemble/_gb.py:494: DataConversionWar
y = column_or_1d(y, warn=True)

```

```
# Prediction
y_train_GB = GB_mod.predict(X_train)
y_train_GB

array([[1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0,
        0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1,
        0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0,
        0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0,
        0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1,
        0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0,
        0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0,
        0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0,
        1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0,
        1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0]])
```

```
# Prediction
y_test_GB = GB_mod.predict(X_test)
y_test_GB

array([[0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0,
        0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0,
        0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1,
        1, 1, 1]])
```

```
print(round(accuracy_score(y_train, y_train_GB), 2))
```

```
1.0
```

```
print(round(accuracy_score(y_test, y_test_GB), 2))
```

```
0.81
```

```
### Confusion Matrix
from sklearn.metrics import confusion_matrix
```

```
confusion_matrix = confusion_matrix(y_test, y_test_GB)
print(confusion_matrix)
```

```
[[33 11]
 [ 6 41]]
```

```
### Classification Report
from sklearn.metrics import classification_report
```

```
print(classification_report(y_test, y_test_GB))
```

```
precision    recall  f1-score   support
```

	0	0.85	0.75	0.80	44
	1	0.79	0.87	0.83	47
accuracy				0.81	91
macro avg		0.82	0.81	0.81	91
weighted avg		0.82	0.81	0.81	91

Knowing your Nearest Neighbours(KNN)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

```
df2=df[:]
```

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
```

```

13 target      303 non-null    int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

from sklearn.preprocessing import LabelEncoder
LE= LabelEncoder()

```

```

df2['sex']= LE.fit_transform(df2['sex'])
df2['cp']= LE.fit_transform(df2['cp'])
df2['fbs']= LE.fit_transform(df2['fbs'])
df2['restecg']= LE.fit_transform(df2['restecg'])
df2['exang']= LE.fit_transform(df2['exang'])
df2['slope']= LE.fit_transform(df2['slope'])
df2['ca']= LE.fit_transform(df2['ca'])
df2['thal']= LE.fit_transform(df2['thal'])
df2

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
0	63	1	1	145	233.0	1	0	150.0	0	2.3	0	0
1	37	1	1	130	250.0	0	1	187.0	0	3.5	0	0
2	41	0	1	130	204.0	0	0	172.0	0	1.4	1	0
3	56	1	1	120	236.0	0	1	178.0	0	0.8	1	0
4	57	0	0	120	354.0	0	1	163.0	1	0.6	1	0
...
298	57	0	0	140	241.0	0	1	123.0	1	0.2	0	0
299	45	1	1	110	264.0	0	1	132.0	0	1.2	0	0
300	68	1	0	144	193.0	1	1	141.0	0	3.4	0	1
301	57	1	0	130	131.0	0	1	115.0	1	1.2	0	1
302	57	0	1	130	236.0	0	0	174.0	0	0.0	0	1

303 rows × 14 columns

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    float64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64

```

```

7   thalach   303 non-null   float64
8   exang     303 non-null   int64
9   oldpeak   303 non-null   float64
10  slope     303 non-null   int64
11  ca        303 non-null   int64
12  thal      303 non-null   int64
13  target    303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

X = df2.loc[:, df2.columns != 'target']
y = df2.loc[:, df2.columns == 'target']

```

```
X.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
dtypes: float64(3), int64(10)
memory usage: 30.9 KB

```

```
y
```

target	
0	1

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
2          1
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```


```
from sklearn.neighbors import KNeighborsClassifier
```

```
290      0
```

```
#Building Model @ n_neighbors = 9
```

```
knn = KNeighborsClassifier(n_neighbors = 9)
print(knn)
mpm_knn = knn.fit(X_train, y_train)
print(mpm_knn)
```

```
KNeighborsClassifier(n_neighbors=9)
KNeighborsClassifier(n_neighbors=9)
/usr/local/lib/python3.7/dist-packages/sklearn/neighbors/_classification.py:198: Data
return self._fit(X, y)
```



```
#Applying on Test data for prediction
y_pred_KNN = mpm_knn.predict(X_test)
print(y_pred_KNN)
```

```
[0 1 1 0 1 1 0 0 0 0 1 1 0 1 1 1 0 1 0 1 1 0 0 0 0 0 1 0 1 1 0 1 1 1 1 0 0
 1 0 1 0 1 1 0 1 1 1 1 1 1 0 1 0 1 0 1 1 1 0 0 1 1 1 0 1 1 1 1 1 1 1 0
 0 0 1 0 1 0 0 0 1 1 0 1 1 0 1 1 1]
```

```
#Prediction Score
mpm_knn.score(X_test, y_test)
```

```
0.7142857142857143
```

```
#Accuracy Score
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred_KNN)
```

```
0.7142857142857143
```

```
# creating a confusion matrix
from sklearn.metrics import confusion_matrix
knn_predictions = knn.predict(X_test)
cm = confusion_matrix(y_test, knn_predictions)
cm
```

```
array([[27, 17],
       [ 9, 38]])
```

Classification Report

```
from sklearn.metrics import classification_report
print(classification_report(y_test, knn_predictions))
```

	precision	recall	f1-score	support
0	0.75	0.61	0.67	44
1	0.69	0.81	0.75	47
accuracy			0.71	91
macro avg	0.72	0.71	0.71	91
weighted avg	0.72	0.71	0.71	91

SVM

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   float64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   float64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB
```

```
df2= df[:]
```

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
```

```

3  trestbps  303 non-null  int64
4  chol      303 non-null  float64
5  fbs       303 non-null  int64
6  restecg   303 non-null  int64
7  thalach   303 non-null  float64
8  exang     303 non-null  int64
9  oldpeak   303 non-null  float64
10 slope     303 non-null  int64
11 ca        303 non-null  int64
12 thal      303 non-null  int64
13 target    303 non-null  int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

from sklearn.preprocessing import LabelEncoder
LE= LabelEncoder()

```

```

df2['sex']= LE.fit_transform(df2['sex'])
df2['cp']= LE.fit_transform(df2['cp'])
df2['fbs']= LE.fit_transform(df2['fbs'])
df2['restecg']= LE.fit_transform(df2['restecg'])
df2['exang']= LE.fit_transform(df2['exang'])
df2['slope']= LE.fit_transform(df2['slope'])
df2['ca']= LE.fit_transform(df2['ca'])
df2['thal']= LE.fit_transform(df2['thal'])
df2

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
0	63	1	1	145	233.0	1	0	150.0	0	2.3	0	0
1	37	1	1	130	250.0	0	1	187.0	0	3.5	0	0
2	41	0	1	130	204.0	0	0	172.0	0	1.4	1	0
3	56	1	1	120	236.0	0	1	178.0	0	0.8	1	0
4	57	0	0	120	354.0	0	1	163.0	1	0.6	1	0
...
298	57	0	0	140	241.0	0	1	123.0	1	0.2	0	0
299	45	1	1	110	264.0	0	1	132.0	0	1.2	0	0
300	68	1	0	144	193.0	1	1	141.0	0	3.4	0	1
301	57	1	0	130	131.0	0	1	115.0	1	1.2	0	1
302	57	0	1	130	236.0	0	0	174.0	0	0.0	0	1

303 rows × 14 columns

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):

```



```

#    Column    Non-Null Count  Dtype
---  -
0    age       303 non-null    int64
1    sex       303 non-null    int64
2    cp        303 non-null    int64
3    trestbps   303 non-null    int64
4    chol       303 non-null    float64
5    fbs        303 non-null    int64
6    restecg    303 non-null    int64
7    thalach    303 non-null    float64
8    exang      303 non-null    int64
9    oldpeak    303 non-null    float64
10   slope      303 non-null    int64
11   ca         303 non-null    int64
12   thal       303 non-null    int64
13   target     303 non-null    int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

X = df2.loc[:, df2.columns != 'target']
y = df2.loc[:, df2.columns == 'target']

```

```
X.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 13 columns):
#    Column    Non-Null Count  Dtype
---  -
0    age       303 non-null    int64
1    sex       303 non-null    int64
2    cp        303 non-null    int64
3    trestbps   303 non-null    int64
4    chol       303 non-null    float64
5    fbs        303 non-null    int64
6    restecg    303 non-null    int64
7    thalach    303 non-null    float64
8    exang      303 non-null    int64
9    oldpeak    303 non-null    float64
10   slope      303 non-null    int64
11   ca         303 non-null    int64
12   thal       303 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.9 KB

```

```
y
```

	target
0	1
1	1
2	1
3	1
4	1
...	...
298	0

```
# splitting
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)

from sklearn import svm
from sklearn.svm import SVC

svm_model= svm.SVC(kernel='linear', C=1, gamma='auto', probability= True).fit(X_train, y_train)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:985: DataConversionWarning:
  y = column_or_1d(y, warn=True)

y_pred_SVM = svm_model.predict(X_test)

# model accuracy for X_test
accuracy = svm_model.score(X_test, y_test)
print(accuracy)

0.8241758241758241

# creating a confusion matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred_SVM)
cm

array([[34, 10],
       [ 6, 41]])


### Classification Report
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_SVM))
```

	precision	recall	f1-score	support
0	0.85	0.77	0.81	44
1	0.80	0.87	0.84	47
accuracy			0.82	91
macro avg	0.83	0.82	0.82	91
weighted avg	0.83	0.82	0.82	91

Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB().fit(X_train, y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:985: DataConversionWarning:
  y = column_or_1d(y, warn=True)
```



```
gnb_predictions = gnb.predict(X_test)
```

```
# accuracy on X_test
accuracy = gnb.score(X_test, y_test)
print(accuracy)
```

```
0.8131868131868132
```

```
# creating a confusion matrix
cm = confusion_matrix(y_test, gnb_predictions)
cm
```

```
array([[35,  9],
       [ 8, 39]])
```

```
### Classification Report
from sklearn.metrics import classification_report
print(classification_report(y_test, gnb_predictions))
```

	precision	recall	f1-score	support
0	0.81	0.80	0.80	44
1	0.81	0.83	0.82	47
accuracy			0.81	91
macro avg	0.81	0.81	0.81	91
weighted avg	0.81	0.81	0.81	91